Hadoop MapReduce Join & Counter with Example

What is a Join in MapReduce?

A join operation is used to combine two large datasets in MapReduce. However, this process involves writing lots of code to perform the actual join operation.

Joining of two datasets begins by comparing the size of each dataset. If one dataset is smaller as compared to the other dataset then smaller dataset is distributed to every data node in the cluster. Once it is distributed, either Mapper or Reducer uses the smaller dataset to perform a lookup for matching records from the large dataset and then combine those records to form output records.

In this tutorial, you will learn-

Types of Join

Depending upon the place where the actual join is performed, this join is classified into-

1. Map-side join - When the join is performed by the mapper, it is called as map-side join. In this type, the join is performed before data is actually consumed by the map function. It is mandatory that the input to each map is in the form of a partition and is in sorted order. Also, there must be an equal number of partitions and it must be sorted by the join key.

2. Reduce-side join - When the join is performed by the reducer, it is called as reduce-side join. There is no necessity in this join to have a dataset in a structured form (or partitioned).

Here, map side processing emits join key and corresponding tuples of both the tables. As an effect of this processing, all the tuples with same join key fall into the same reducer which then joins the records with same join key.

An overall process flow is depicted in below diagram.

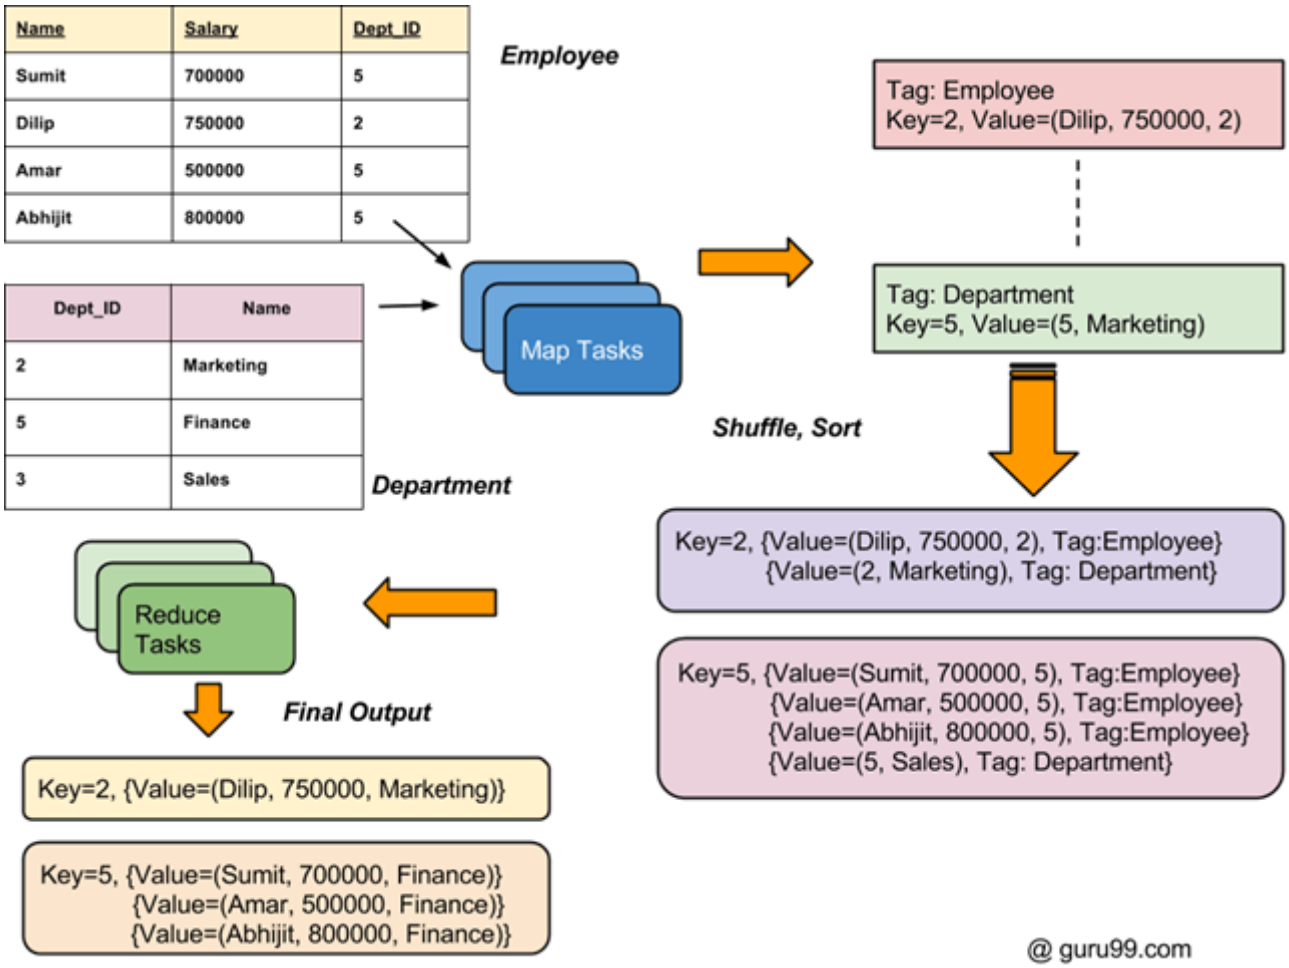

How to Join two DataSets: MapReduce Example

There are two Sets of Data in two Different Files (shown below). The Key Dept_ID is common in both files. The goal is to use MapReduce Join to combine these files

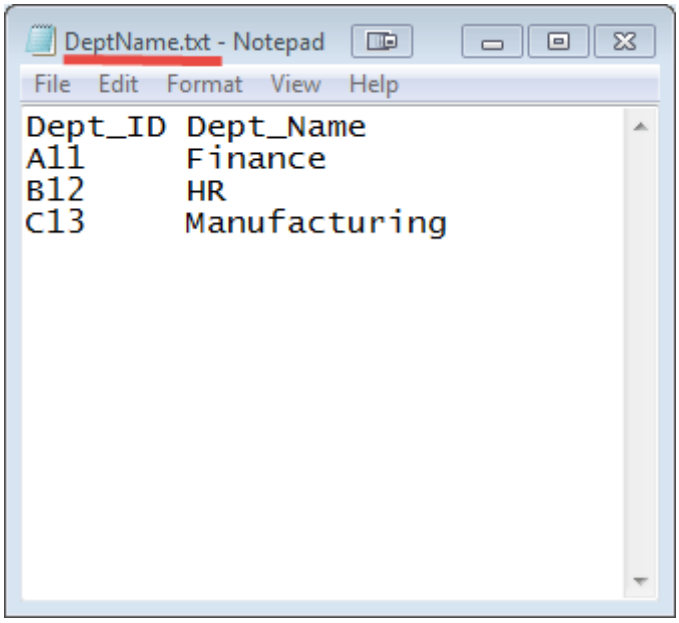

File 1

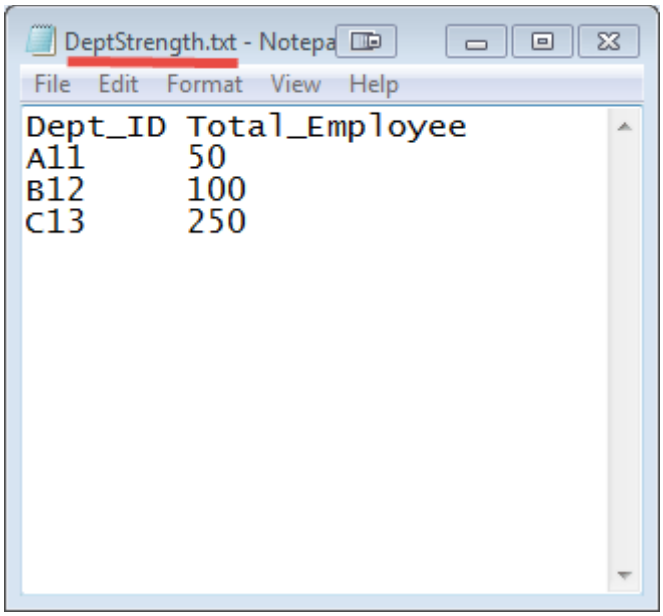

File 2

Input: The input data set is a txt file, **DeptName.txt & DepStrength.txt**

Download Input Files From Here

Ensure you have Hadoop installed. Before you start with the actual process, change user to 'hduser' (id used while Hadoop configuration, you can switch to the userid used during your Hadoop config).

```
su - hduser_
```

```
guru99@guru99-VirtualBox:~$ su - hduser_
Password:
hduser_@guru99-VirtualBox:~$
```

Step 1) Copy the zip file to the location of your choice

```
hduser_@guru99-VirtualBox:~$ cp  /home/guru99/Downloads/MapReduceJoin.tar.gz /ho
me/hduser_/
hduser_@guru99-VirtualBox:~$ ls /home/hduser_
apache-flume-1.4.0-bin  guava-17.0.jar        MapReduceTutorial
examples.desktop        hdfs                  pig_1399372687264.log
FlumeTutorial           inputMapReduce        protobuf-java-2.4.1.jar
guava-10.0.1.jar        MapReduceJoin.tar.gz
```

Step 2) Uncompress the Zip File

```
sudo tar -xvf MapReduceJoin.tar.gz
```

```
hduser_@guru99-VirtualBox:~$ sudo tar -xvf MapReduceJoin.tar.gz
[sudo] password for hduser_:
MapReduceJoin/
MapReduceJoin/TextPair.java
MapReduceJoin/MapReduceJoin.jar
MapReduceJoin/JoinReducer.java~
MapReduceJoin/Manifest.txt
MapReduceJoin/DeptEmpStrengthMapper.java~
MapReduceJoin/JoinReducer.java
MapReduceJoin/TextPair.java~
MapReduceJoin/DeptNameMapper.java~
MapReduceJoin/JoinDriver.java
MapReduceJoin/Manifest.txt~
MapReduceJoin/DeptNameMapper.java
MapReduceJoin/DeptEmpStrength.txt
MapReduceJoin/JoinDriver.java~
MapReduceJoin/A.txt~
MapReduceJoin/B.txt~
MapReduceJoin/MapReduceJoin/
MapReduceJoin/MapReduceJoin/TextPair$FirstComparator.class
MapReduceJoin/MapReduceJoin/DeptNameMapper.class
MapReduceJoin/MapReduceJoin/JoinDriver$KeyPartitioner.class
MapReduceJoin/MapReduceJoin/TextPair.class
MapReduceJoin/MapReduceJoin/JoinDriver.class
MapReduceJoin/MapReduceJoin/TextPair$Comparator.class
MapReduceJoin/MapReduceJoin/JoinReducer.class
MapReduceJoin/MapReduceJoin/DeptEmpStrengthMapper.class
MapReduceJoin/DeptEmpStrengthMapper.java
MapReduceJoin/DeptName.txt
MapReduceJoin/DeptStrength.txt
hduser_@guru99-VirtualBox:~$
```

Step 3) Go to directory MapReduceJoin/

```
cd MapReduceJoin/
```

```
hduser_@guru99-VirtualBox:~$ cd MapReduceJoin
hduser_@guru99-VirtualBox:~/MapReduceJoin$
```

Step 4) Start Hadoop

```
$HADOOP_HOME/sbin/start-dfs.sh
$HADOOP_HOME/sbin/start-yarn.sh
```

```
hduser_@guru99-VirtualBox:~/MapReduceJoin$  $HADOOP_HOME/sbin/start-dfs.sh
Starting namenodes on [localhost]
localhost: starting namenode, logging to /home/guru99/Downloads/hadoop/logs/hadoop-hduser_-namenode-guru99
-VirtualBox.out
localhost: starting datanode, logging to /home/guru99/Downloads/hadoop/logs/hadoop-hduser_-datanode-guru99
-VirtualBox.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /home/guru99/Downloads/hadoop/logs/hadoop-hduser_-secondar
ynamenode-guru99-VirtualBox.out
hduser_@guru99-VirtualBox:~/MapReduceJoin$  $HADOOP_HOME/sbin/start-yarn.sh
starting yarn daemons
starting resourcemanager, logging to /home/guru99/Downloads/hadoop/logs/yarn-hduser_-resourcemanager-guru9
9-VirtualBox.out
localhost: starting nodemanager, logging to /home/guru99/Downloads/hadoop/logs/yarn-hduser_-nodemanager-gu
ru99-VirtualBox.out
hduser_@guru99-VirtualBox:~/MapReduceJoin$ $HADOOP_HOME/bin/hdfs dfs -copyFromLocal DeptStrength.txt DeptN
ame.txt /
hduser_@guru99-VirtualBox:~/MapReduceJoin$
```

Step 5) DeptStrength.txt and DeptName.txt are the input files used for this program.

These file needs to be copied to HDFS using below command-

```
$HADOOP_HOME/bin/hdfs dfs -copyFromLocal DeptStrength.txt DeptName.txt /
```

```
hduser_@guru99-VirtualBox:~/MapReduceJoin$ $HADOOP_HOME/bin/hdfs dfs -copyFromLocal DeptStrength.txt DeptN
ame.txt /
hduser_@guru99-VirtualBox:~/MapReduceJoin$
```

Step 6) Run the program using below command-

```
$HADOOP_HOME/bin/hadoop jar MapReduceJoin.jar MapReduceJoin/JoinDriver/DeptStrengt
h.txt /DeptName.txt /output_mapreducejoin
```

```
hduser_@guru99-VirtualBox:~/MapReduceJoin$ $HADOOP_HOME/bin/hadoop jar MapReduceJoin.jar /DeptStrength.txt
 /DeptName.txt /output_mapreducejoin
```

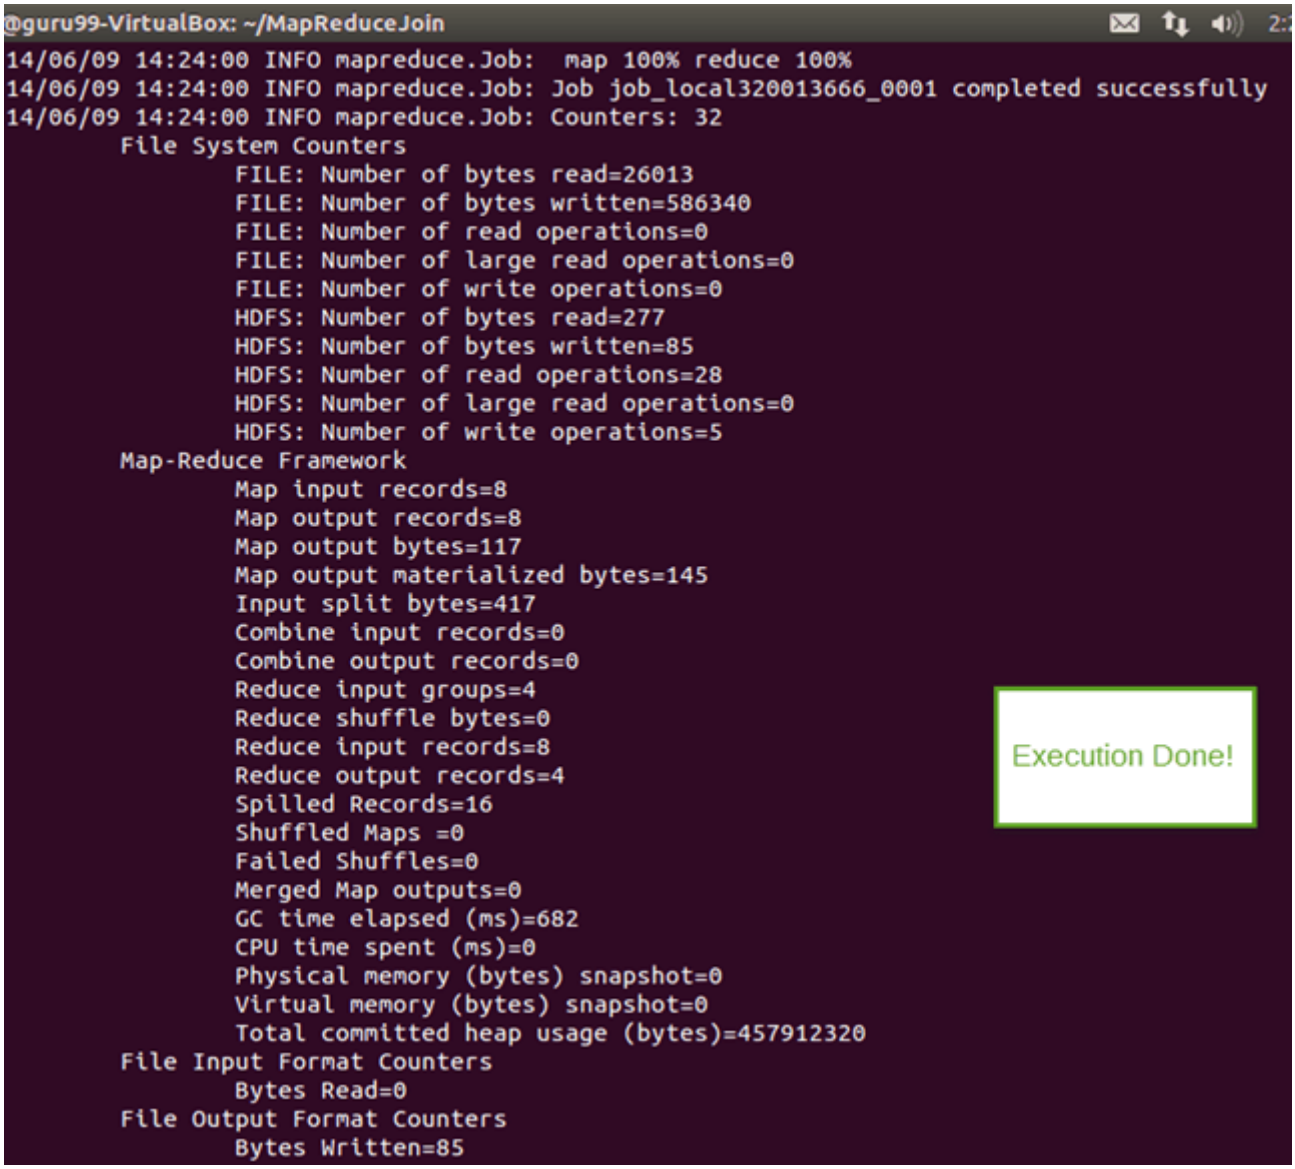

Step 7) After execution, output file (named 'part-00000') will stored in the directory /output_mapreducejoin on HDFS

Results can be seen using the command line interface

```
$HADOOP_HOME/bin/hdfs dfs -cat /output_mapreducejoin/part-00000
```

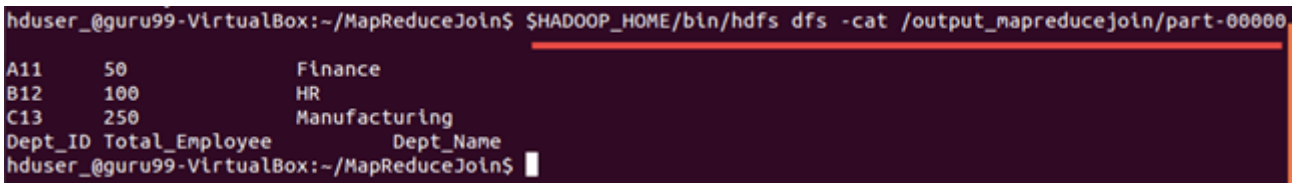

Results can also be seen via a web interface as-

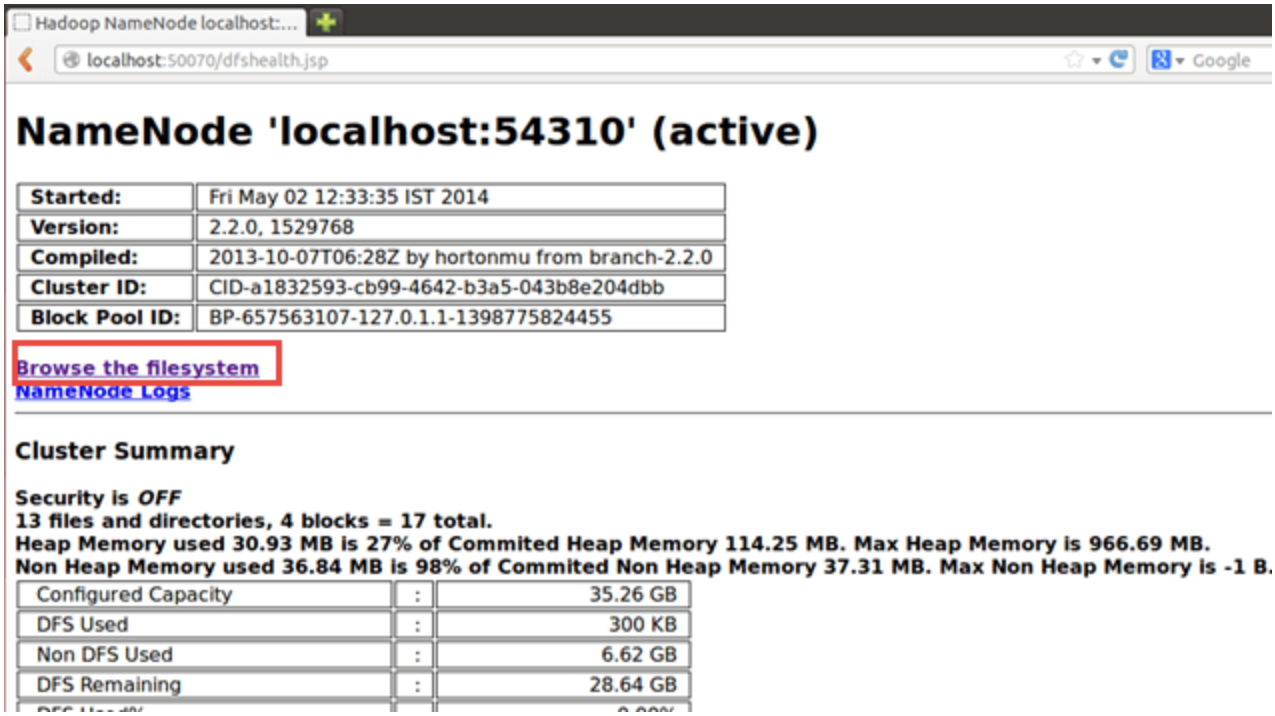

Now select **'Browse the filesystem'** and navigate upto **/output_mapreducejoin**

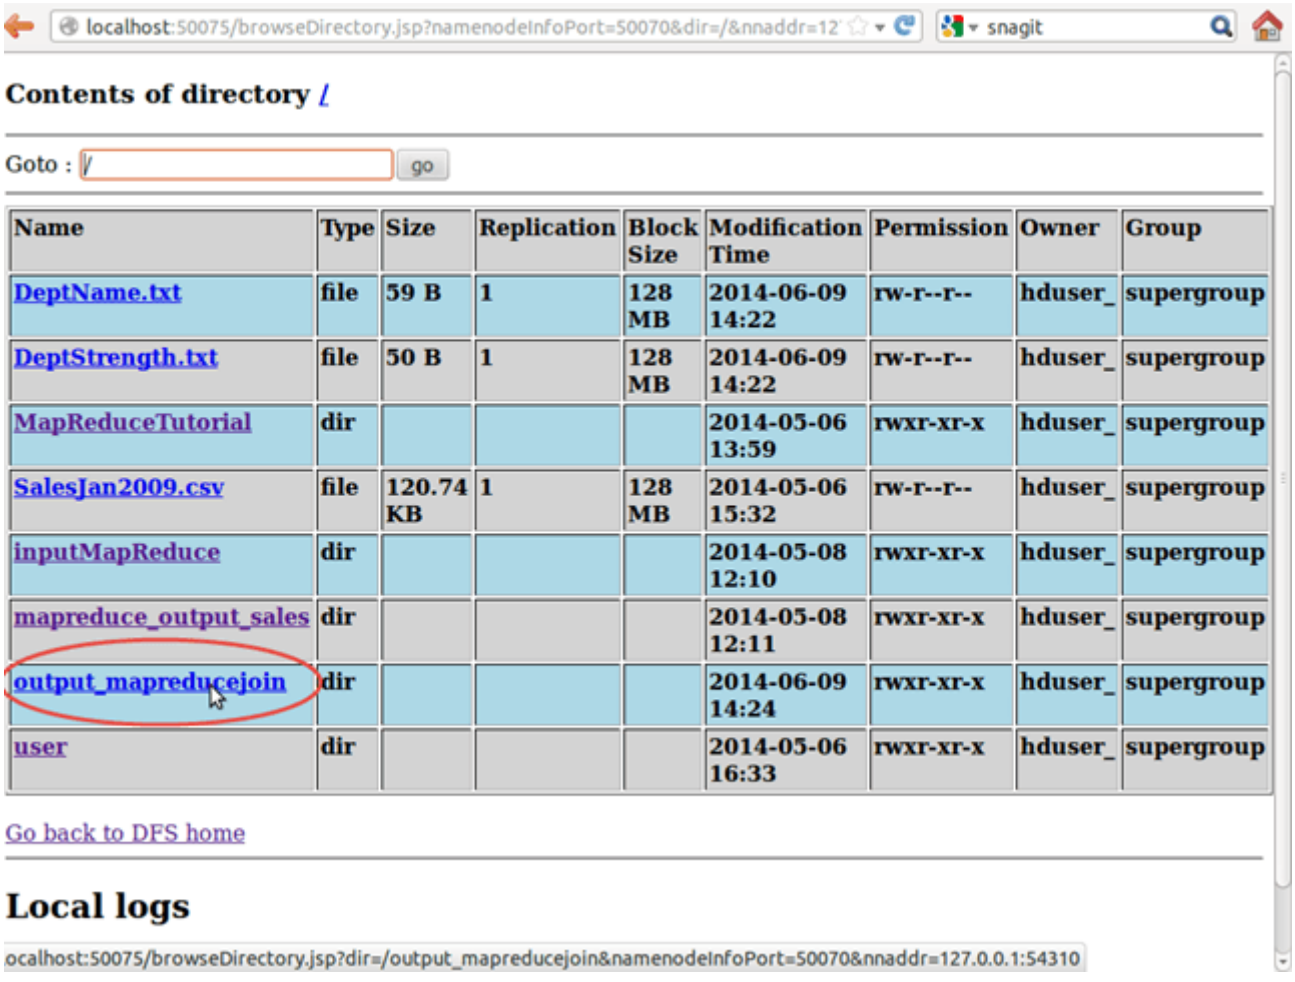

Open **part-r-00000**

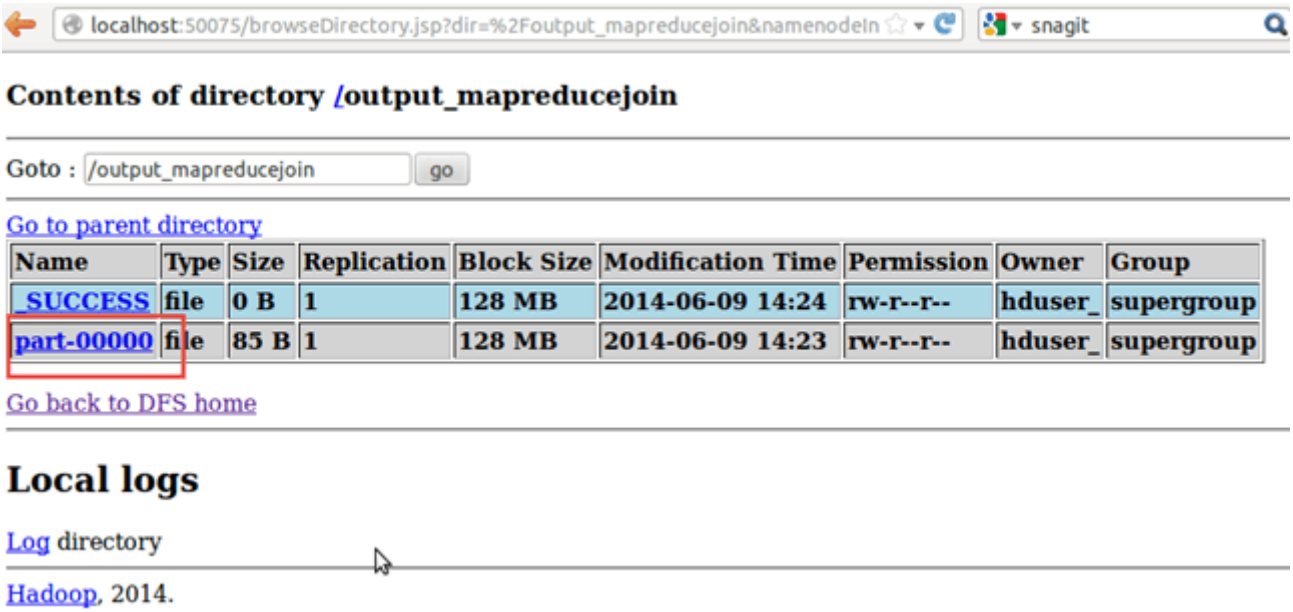

Results are shown

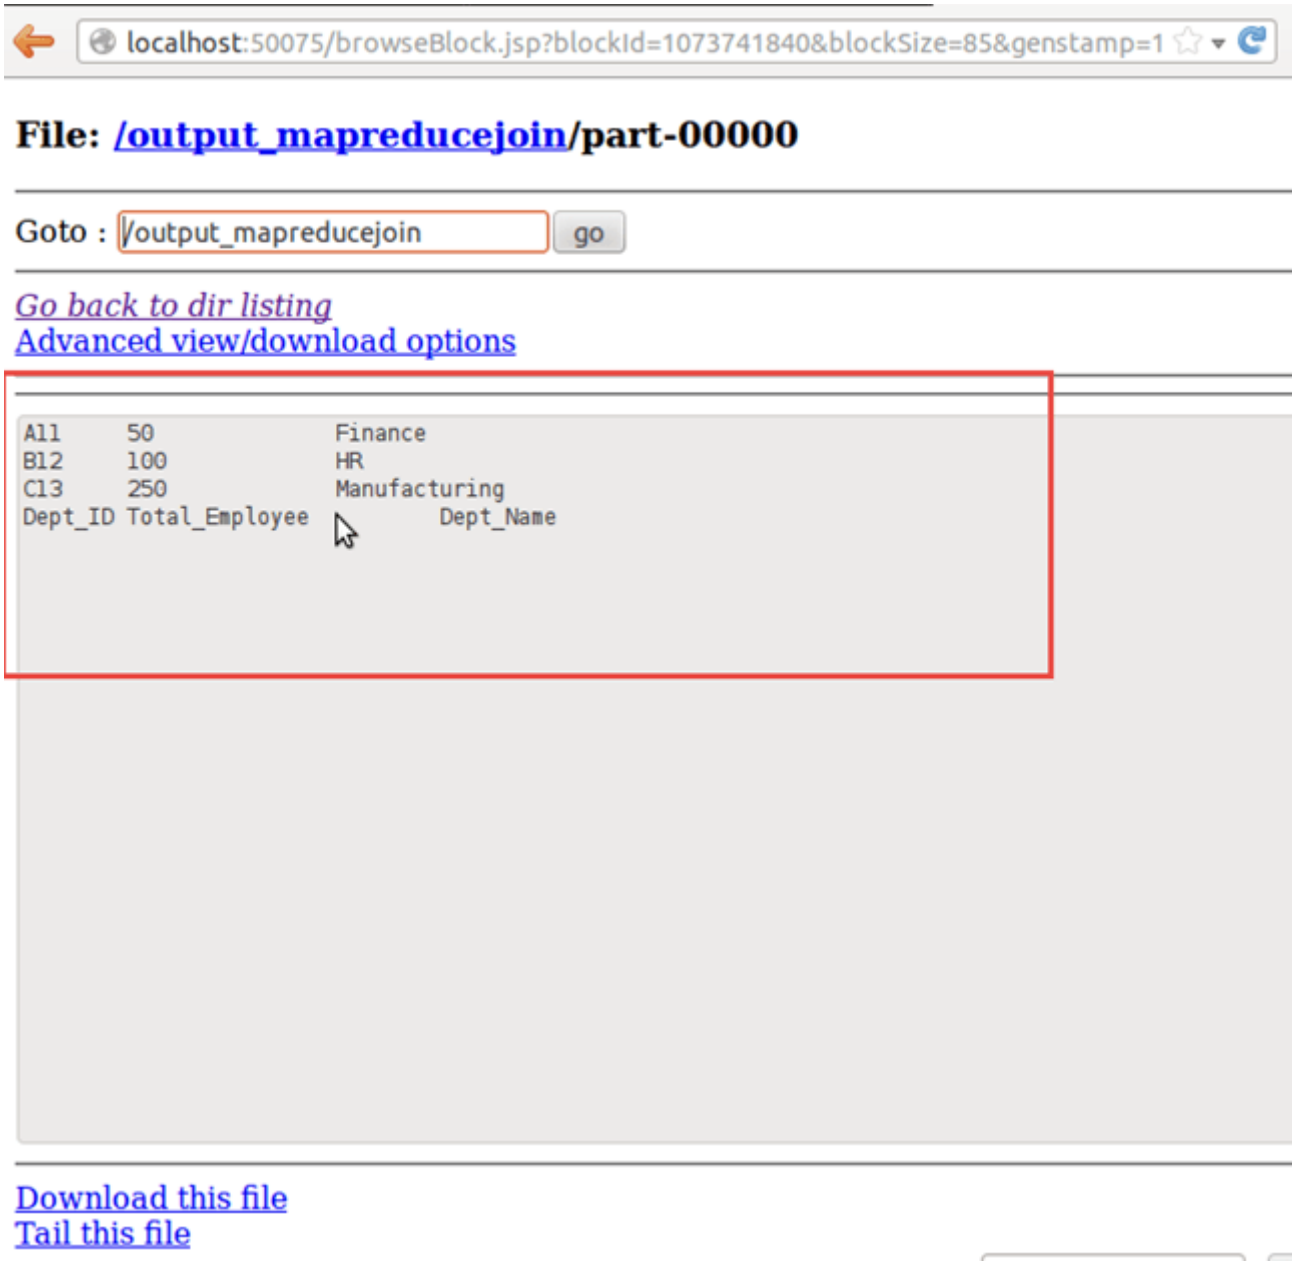

NOTE: Please note that before running this program for the next time, you will need to delete output directory /output_mapreducejoin

```
$HADOOP_HOME/bin/hdfs dfs -rm -r /output_mapreducejoin
```

Alternative is to use a different name for the output directory.